

# CP322 Project

Group 9:

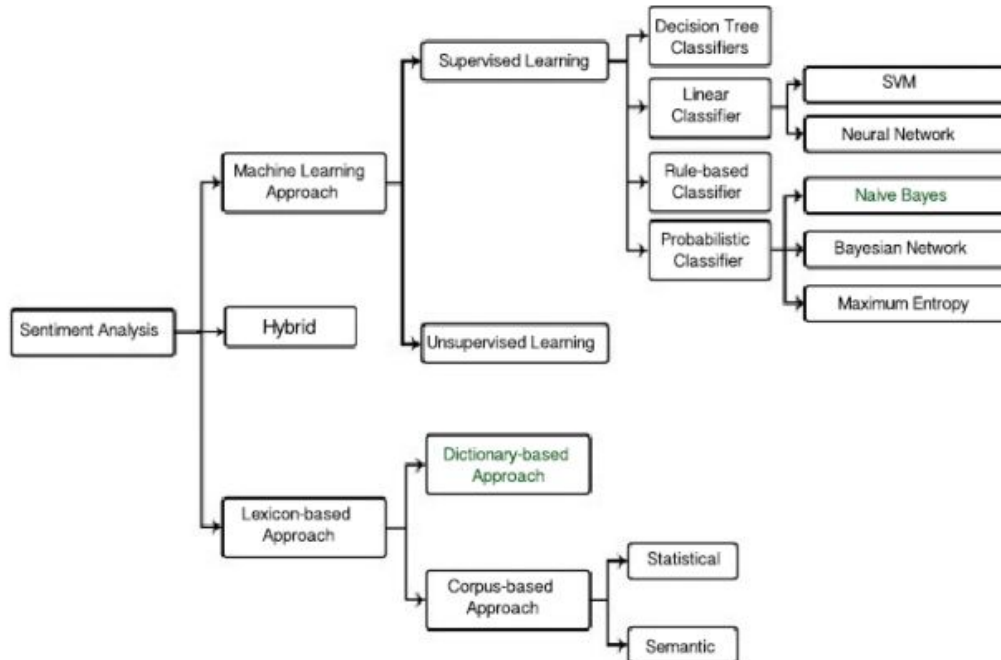
Phoebe Schulman, Anojan Balendra, Harchit Bhatoia,  
Shing Hei Mau, Alvin Onabolu

# Introduction

	A	B	C	D	E	F	G	H	I	J	K	L
1		Clothing ID	Age	Title	Review Text	Rating	Recommendation	Positive Feedback	Division Name	Department	Class Name	
2	0	767	33	Absolutely		4	1	0	Initmates	Intimate	Intimates	
3	1	1080	34	Love this di		5	1	4	General	Dresses	Dresses	
4	2	1077	60	Some major had such l		3	0	0	General	Dresses	Dresses	
5	3	1049	50	My favorite! I love, love,		5	1	0	General Pe	Bottoms	Pants	
6	4	847	47	Flattering s This shirt is		5	1	6	General	Tops	Blouses	
7	5	1080	49	Not for the l love tracy		2	0	4	General	Dresses	Dresses	
8	6	858	39	Cagrcosal I aded this		5	1	1	General Pe	Tops	Knits	
9	7	858	39	Shimmer, s I ordered tl		4	1	4	General Pe	Tops	Knits	
10	8	1077	24	Flattering I love this c		5	1	0	General	Dresses	Dresses	
11	9	1077	34	Such a fun I'm 5"5' an		5	1	0	General	Dresses	Dresses	
12	10	1077	53	Dress look: Dress runs		3	0	14	General	Dresses	Dresses	

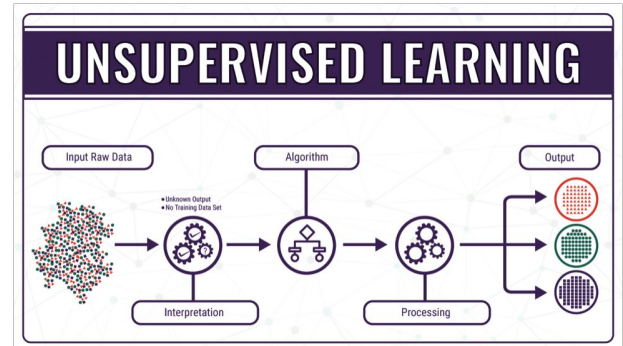
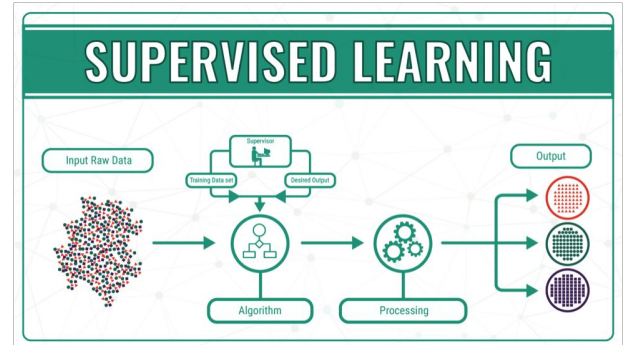
- Given data set of clothing reviews
- **Question being addressed**
  - How to develop a sentiment classifier that assigns reviews a rating from 1-5?
- Sentiment analysis
  - Analyzing text
  - Determine if an opinion is positive or negative

# Related Work: Sentiment Analysis



# Related Work: Machine Learning

- Supervised or Unsupervised?
- Naive Bayes
- K-Nearest Neighbors
- Support Vector Machine (SVM)
- Decision Trees



# Naive Bayes + TFIDF

TFIDF importance score

```
{'major': 1986,
 'design': 918,
 'flaw': 1275,
 'favorite': 1207,
 'buy': 493,
 'flattering': 1270,
 'shirt': 2831,
 'petite': 2364,
 'cagrcoal': 504,
 'shimmer': 2824,
 'fun': 1365,
 'surprisingly': 3175,
 'go': 1428,
 'lot': 1924,
 'dress': 1000,
 'look': 1901,
 'like': 1867,
 'make': 1989,
 'cheap': 578,
```

Word's Unique  
value

index

```
(2, 918) 1
(2, 1275) 1
(2, 1986) 1
(3, 493) 1
(3, 1207) 1
(4, 1270) 1
(4, 2831) 1
(5, 2364) 1
(6, 504) 1
(6, 1365) 1
(6, 2824) 1
(7, 1428) 1
(7, 1924) 1
(7, 2824) 1
(7, 3175) 1
(8, 1270) 1
(9, 1000) 1
(9, 1365) 1
(10, 578) 1
(10, 1901) 1
```

Occurrence  
frequency

```
(2, 1986) 0.7052716454501263
(2, 1275) 0.5691706469378793
(2, 918) 0.42265432777671547
(3, 1207) 0.6725653380818224
(3, 493) 0.7400377463419581
(4, 2831) 0.6619290578025505
(4, 1270) 0.7495664896702813
(5, 2364) 1.0
(6, 2824) 0.616833907181881
(6, 1365) 0.3567917478968672
(6, 504) 0.7015807719596033
(7, 3175) 0.4977749905370736
(7, 2824) 0.5923078260116152
(7, 1924) 0.479087681657613
(7, 1428) 0.41456783681943293
(8, 1270) 1.0
(9, 1365) 0.8220129508290713
```

```
bow = transformer.transform(df["Title"])
print(bow)
```

```
title_tfidf = tfidf_trans.transform(bow)
print(title_tfidf)
```

# Text data

vs

# title data

```
array([[ 0,  0, 22, 23, 797],  
       [ 0,  3, 29, 39, 1494],  
       [ 0,  0, 53, 40, 2778],  
       [ 0,  0,  1, 148, 4928],  
       [ 0,  0,  0,  4, 13127]], dtype=int64)
```

```
array([[ 35,  61, 243,  89, 414],  
       [  7, 151, 439, 262, 706],  
       [  0,  25, 928, 499, 1419],  
       [  0,  3, 128, 1179, 3767],  
       [  0,  2,  28, 207, 12894]], dtype=int64)
```

Overall Accuracy: 64.66%

Overall Accuracy: 56.76%

Class	n (truth) ②	n (classified) ②	Accuracy	Precision	Recall	F1 Score
1	0	842	96.41%	0.0	0.0	0.0
2	3	1565	93.35%	0.0019	1.0	0.0038
3	105	2871	87.78%	0.018	0.50	0.036
4	254	5077	78.56%	0.029	0.58	0.056
5	23124	13131	57.42%	1.0	0.57	0.72

Class	n (truth) ②	n (classified) ②	Accuracy	Precision	Recall	F1 Score
1	42	842	96.53%	0.042	0.83	0.079
2	242	1565	93.59%	0.096	0.62	0.17
3	1766	2871	88.16%	0.32	0.53	0.40
4	2236	5077	78.9%	0.23	0.53	0.32
5	19200	13131	72.14%	0.98	0.67	0.80

# Solution/Method

## What was done

- Sentiment classification during preprocessing
- Predictive classifier: K-Nearest Neighbors
  - Why: better results than Gaussian Naive Bayes
  - Where K value changes
    - Why: best K value changed each time

```
# find a k value for knn- Look for Lowest point(error) on graph

minError = 1 #where is least error
minErrorK = 0 #which k to choose
errors = [0] #errors = []

# Calculate error for dif K values
for i in range(1,35): #(5,35): #(1, 40)

    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(descriptive_train, target_train)
    pred_i = knn.predict(descriptive_test)
    errors.append(np.mean(pred_i != target_test))

    #print(errors[i-1], np.mean(pred_i != target_test), minError, minErrorK)
    if(errors[i] < minError): #if(errors[i-1] < minError):
        minError = errors[i]
        minErrorK = i

errors = errors[1:]
print("choose k = ", minErrorK)


plt.figure(figsize=(12, 6))
plt.plot(range(1, 35), errors, color='red', linestyle='dashed', marker='o', markerfacecolor='blue',

plt.title('Error Rate K Value')
plt.xlabel('K Value')
plt.ylabel('Mean Error')
```

# Solution/Method

## Tools and techniques

- Finding best K value for KNN
  - Run many KNN
  - Calculate error for each K value
  - Find lowest point (least error) on graph
  - Use the corresponding K for actual predictions



choose k = 27



# Solution/Method

## **Tools and techniques**

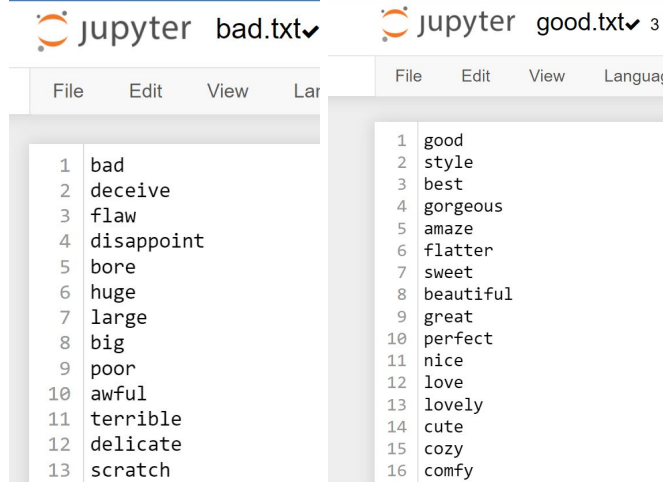
- Tokenization
  - Why: convert sentences into words
- Lemmatization
  - Reduces words to their root word
  - Eg: flattering -> flatter, flattered -> flatter
  - Why: not need to consider every version of a word



# Data and Experiments

## Preprocessing data

- Read files containing good or bad words
  - Added specific words from data
    - Eg: chic, cute, fab, itchy
- Sentiment analysis on Title and Review Text



The image shows two side-by-side Jupyter Notebook windows. The left window is titled 'bad.txt' and contains a list of 13 negative words. The right window is titled 'good.txt' and contains a list of 16 positive words. Both windows have a menu bar with 'File', 'Edit', 'View', and 'Language' options.

bad.txt	good.txt
1 bad	1 good
2 deceive	2 style
3 flaw	3 best
4 disappoint	4 gorgeous
5 bore	5 amaze
6 huge	6 flatter
7 large	7 sweet
8 big	8 beautiful
9 poor	9 great
10 awful	10 perfect
11 terrible	11 nice
12 delicate	12 love
13 scratch	13 lovely
	14 cute
	15 cozy
	16 comfy

# Data and Experiments

## Preprocessing data

- Evaluating Title and Review Text features
  - Replace nan with 0
  - Finding number of good/bad words in sentences
  - **final\_eval = num\_good\_words - num\_bad\_words**
    - Positive when more good words
    - Negative when more bad words
  - Handle negation (“no”, “not”) -> switch the outcome
  - Handle expression (“!” , “too”) -> increase the outcome

```
#title
title_evals= []

for i in range(len(sentiment_df_title)):#for each sentence
    num_good_words = 0
    num_bad_words = 0

    final_eval = 0

    was_negated = False
    was_expression = False

    if (sentiment_df_title.values[i]!=0):#not na

        sentence = sentiment_df_title.values[i]
        print(sentence)#debug
        words=word_tokenize(sentence)

        for j in range(len(words)): #for each word

            word = words[j].lower()
            rootWord = lem.lemmatize(word,"v") #root words

            #meaning of word
            if rootWord in goodWords :
                num_good_words+=1

            elif rootWord in badWords:
                num_bad_words+=1

            elif rootWord in negateWords:
                was_negated = True
            elif rootWord in expressionWords:
                was_expression = True
```

# Data and Experiments

## Examples

- Bad: Itchy
- Negated: Not impressed
  - impressed-> impress which is good
  - Not added -> bad
- Expression: Great shirt!!!
  - Great is positive
  - ! -> even more positive
- Combination: Not very flattering

```
Itchy tags
was_negated, was_expression =  False False
num_good_words, num_bad_words =  0 1
final_eval = -1
```

```
Not impressed...
was_negated, was_expression =  True False
num_good_words, num_bad_words =  1 0
final_eval = -1
```

```
Great shirt!!!
was_negated, was_expression =  False True
num_good_words, num_bad_words =  1 0
final_eval = 2
```

```
Not very flattering
was_negated, was_expression =  True True
num_good_words, num_bad_words =  1 0
final_eval = -2
```



descriptive\_features=====

# Data and Experiments

	Recommended IND	Positive Feedback Count	titleNew	textNew
0	1	0	0	3
1	1	4	0	8
2	0	0	-1	6
3	1	0	2	-14
4	1	6	1	6

## Preprocessing data

- Creating descriptive vs target features
  - Target features: Rating
  - Descriptive features: Title, Review Text, Positive Feedback Count, Recommended IND
    - Dropping remaining features
    - Replace Title and Review Text with final evaluations
- Randomly split original data
  - 70% training : 30% testing



# Data and Experiments

## Reliable

- Robustness of good and bad words approach largely depends on the supply of good/bad words being checked against
- Larger word bank -> better results -> more reliable



# Evaluation and Results

F1score: 0.555832

recall: 0.619075

## Model evaluations

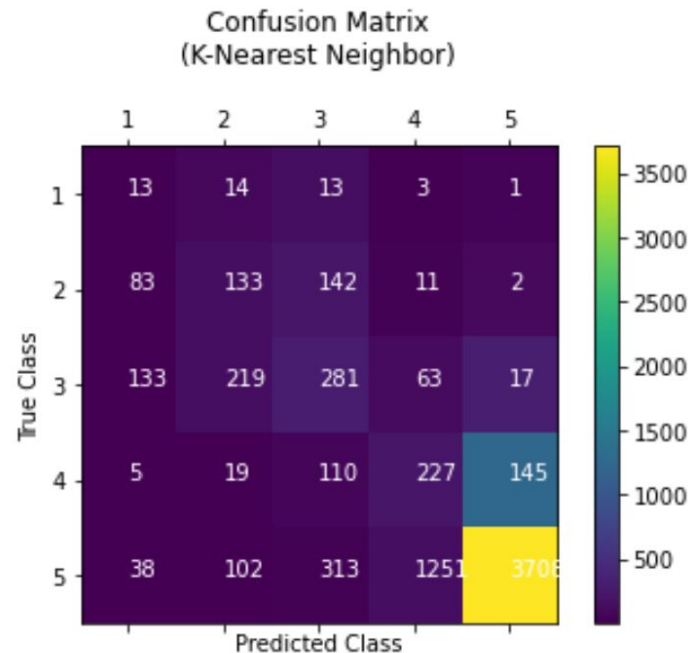

- F1 score: 55%
- Recall: 62%
- Others
  - F1 score for each rating (1-5)
  - Accuracy: 62%

	precision	recall	f1-score	support
1	0.05	0.30	0.08	44
2	0.27	0.36	0.31	371
3	0.33	0.39	0.36	713
4	0.15	0.45	0.22	506
5	0.96	0.69	0.80	5412
accuracy			0.62	7046
macro avg	0.35	0.44	0.35	7046
weighted avg	0.79	0.62	0.68	7046

# Evaluation and Results

## Model evaluations

- Confusion matrix shows final ratings
- Eg:
  - 1,1 -> true positive of rating 1 (13)
  - 5,5 -> true positive of rating 5 (3708)
  - Why: imbalance data





# Evaluation and Results

## Conclusions

- Best results seen in KNN approach
- More robustness with manual preprocessing of data

```
IN CONCLUSIONS...we get results similar to :  
  
F1score: 0.555832  
precision: 0.559843  
recall: 0.619075  
  
      precision    recall  f1-score   support  
  
     1      0.05      0.30      0.08        44  
     2      0.27      0.36      0.31       371  
     3      0.33      0.39      0.36       713  
     4      0.15      0.45      0.22       506  
     5      0.96      0.69      0.80      5412  
  
   accuracy      0.62      7046  
  macro avg      0.35      0.44      0.35      7046  
 weighted avg      0.79      0.62      0.68      7046  
  
therefore our model is about 62% accurate, F1score is 56%, and recall is 62% .  
  
In addition:  
f1 score of each target is different.  
eg:  
rating 1 with f1score= 0.08  
vs  
rating 5 with f1score= 0.80  
  
this is either because the data is highly imbalanced or because we are better at predicting the 5 star ratings.  
...
```

# Q and A

